

High-Performance Computing (Resources and Capabilities)

<https://github.com/wsphd/csun-hpc/>

"High-Performance Computing (Resources and Capabilities)
California State University, Northridge (CSUN)

Friday, April 11, 2025
CSUN Faculty Retreat - Odyssey Restaurant

Wayne Smith, Ph.D., *Department of Management*, ws@csun.edu

(Note: This page is also available in alternate formats - [.pdf](#), [.odt](#), and [.docx](#))

Introduction/Background/Motivation

- Some n i N needs are \leq contemporary desktop/laptop and software
 - But double-check new methodologies and growth (and by extension, movement) of data
- Some n j N needs are $>$ contemporary desktop/laptop and software
 - Essentially, "compute-intensive, data-intensive, or network-intensive"
 - Use primarily FOSS (Linux, Open Source, etc.) to complement COTS (Windows, SPSS, etc.)
- Private, "on-premises" servers (contact: [Zack Hillbruner](#), CSUN IT)
 - Usually purchased by an individual faculty member or Dept. (often with a grant or project)
 - Usually located in the on-campus CSUN MDF
 - CSUN IT usually racks and networks the system; Users manage the system and applications
- Public Cloud (AWS, GCP, MS-Azure, OCI, Digital Ocean, etc.)
 - Use "free-tier" (still need to provide a credit card)
 - Purchase credits with a credit card
 - Scholars can ask for resources for substantive research
- Or?
 - NSF-funded, multi-year, inter-institution, STEAM/SocialSTEM, R3s/CCC's too
 - [CSUN IT Technology Resources for Research](#)

General Advanced Computing/Data Management

- There are plenty of (non-HPC) advanced computing issues too (research and instruction).
- Ecosystem Transition: Compute

- Beyond CPUs, there GPUs, FPGAs, DPUs, and others
 - COTS languages (e.g., SPSS, Stata, MPlus, Matlab, NVivo) -> FOSS languages (e.g., R, Python, Julia)
 - Single-threaded execution -> Parallel execution
 - COTS spreadsheets (e.g., Excel) -> FOSS spreadsheets(e.g., LibreOffice).
 - Operating Systems (e.g., Windows/MacOs) *plus* Linux, Excel -> LibreOffice, etc.
 - Beyond replication -> Reproducibility (not just 'A' journals)
- Ecosystem Transition: Data
 - "Big Data"
 - Research results can include output data (and perhaps even source data) too
 - Desktop sizes (e.g., GiB, TiB) -> Beyond-Desktop sizes (e.g., PiB, EiB)
 - Human-readable file formats with no meta-data (e.g., CSV) -> Digital file formats with meta-data (e.g., Parquet)
 - Row-oriented databases (e.g., MariaSQL, PostgreSQL, SQLite, etc.) -> Column-oriented databases (e.g., DuckDB, MonetDB, TileDB)
 - Monolith APIs (e.g., REST) -> Robust APIs (e.g., GraphQL)
 - Single-file access (e.g., HTTPS) -> Multi-file 'buckets' (e.g., AWS/S3, GCS, Azure/Blob)
- Ecosystem Transition: Network
 - "High Throughput"
 - Big Data needs to be moved over fast, reliable networks
 - CSUN 'Science DMZ'
 - Los Angeles Public Library (LAPL) Northridge branch and Mid-Valley branch (>2x file transfer, symmetric)
- Example: Technology Trends
 - Campus Labs *plus* Home Labs, Open Science, Open Research, Open Data, Open anything...
- I'm happy to discuss these issues too but it's not the primary focus of this material.

Faculty/Disciplines

- CSUN is a big place--there are many faculty doing interesting things with HPC
- Ravi Absol (Chemistry)
 - NRP
 - Molecular Dynamics
 - undergraduate student (Anita) - uses AMBER software (with many GPUs, including NVIDIA A100s)
 - one protein combinatorics imulation took 3 months on a laptop -- now takes ~ 3 weeks
- Bingbing Li (Manufacturing Systems Engineering and Management)
 - NRP
 - Smart Manufacturing (e.g., Industrial AI, Multimodal Data Fusion for Autonomous System, Digital Twins, XR and Metaverse for Manufacturing)
- Xiyi Hang (Electrical and Computer Engineering)
 - NRP

- Bioinformatics/Bio-medical applications (e.g., Gene expression, data mining)
- Xunfei Jiang (Computer Science)
 - on-premises HPC and OSG/NRP
 - Cloud and Infrastructure design and management
 - Workforce development (full stack)
- Dr. Nicholas Kioussis (Physics)
 - CalTech Quantum Computing (and others)
 - Materials Science
- Wayne Smith (Management)
 - ACCESS and NRP
 - Federal Communications Commission (FCC) Universal Licensing System (ULS) data
 - Github - [R code](#)
 - Data output - [multi-file data output](#)
- Other future CSUN projects
 - CTVA - UNREAL 3D imaging (e.g., digital twins, gaming, 3D mapping)
 - Digital Humanities
- Representative CSU projects
 - SDSU
 - CSUSB - [High Performance Computing](#)
 - Sonoma State
 - CalPolyHumboldt
- CSU help
 - Dr. Dung Vu (CSUSB)
 - Kyle Krick (SDSU)

JupyterHub

- [Multi-user, interactive notebooks (including R, Python, Julia, etc.)]<https://jupyter.org/hub>
- CSU Technology Infrastructure for Exploration ("TIDE")
 - *governance*: funded by NSF, provisioned by Cloudbank, managed/supported by UC Berkeley
 - *hardware*: basic access
 - *software*: R/RStudio, Python, Shiny, Linux terminal, generic notebooks
 - *log in*: Select "California State University, Northridge" and then log in is via CILOGON (just log in as usual)
 - Start here: [CAL-ICOR]<https://csun.cloudbank.2i2c.cloud/>
- CSU Technology Infrastructure for Exploration ("TIDE")
 - *governance*: funded by NSF, provisioned by NRP, managed/supported by SDSU
 - *hardware*: can have access to multiple GPUs, multiple cores, lots of RAM, some (shareable) storage
 - *software*: R/RStudio, Tensorflow, Datascience, generic notebooks, discipline-specific notebooks
 - *log in*: log in is via CILOGON (just select "California State University, Northridge" and log in as usual)

- Start here: [CSUN TIDE]<https://tide.sdsu.edu/>
- CSUN myCSUNSoftware (aka "Apporto")
 - n/a (but other GUI and CLI software runs here)
 - Faculty will need to ask for it over time
 - Start here: [myCSUNSoftware]<https://www.csun.edu/it/software-services/software/all-software/mycsunsoftware>

Jetstream2/ACCESS

- Managed by Indiana University
- 100's of GiB of RAM, 10's of PB of disk, 10's of [GPUs](#), fast networks
 - Best for new learners, data science projects (R, Python, Julia, etc.), large simulations, gateway to other systems, including several supercomputers around the country
- Need an "ACCESS ID"
 - Like an ORCID ID but for Research Computing
 - Have CV or Resume for upload (don't worry, your request will be approved)
- Be willing to learn:
 - How to ask (nicely and well, for more (incrementally) resources, and read a simple dashboard
 - the Command line and Linux
 - Webshell
 - SSH for logging int (and some learning curve for generating SSH keys and passphrases)
 - SCP for file transfer (after the SSH process is done)
- (Live demo...if possible)
- Start here:
 - [Jetstream ACCESS page]<https://jetstream-cloud.org/get-started/index.html>

Nautilus/NRP

- Managed by University of California, San Diego
- 100's of GiB of RAM, 10's of PB of disk, 100's of [GPUs](#)/[FPGAs](#)/[TPUs](#)/[DPUs](#), very fast networks
 - Best for leading-edge science and engineering, especially w/ funded labs and staff
- Be willing to learn:
 - Must be comfortable with the Command Line, Open Source, and Linux
 - Kubernetes (open source client-server), you use the "kubectl" binary
 - You control just about everything with ASCII ".yaml" files
- (Static demo -- sample .yaml configuration file with GPUs)

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deployment-amber-gpu-ws
  labels:
```

```

    k8s-app: deployment-amber-gpu-ws
spec:
  replicas: 1
  selector:
    matchLabels:
      k8s-app: deployment-amber-gpu-ws
  template:
    metadata:
      labels:
        k8s-app: deployment-amber-gpu-ws
    spec:
      containers:
        - name: mypod-gpu
          image: gitlab-registry.nrp-nautilus.io/prp/jupyter-stack/tensorflow
          resources:
            limits:
              memory: 512Gi
              cpu: 1500m
              nvidia.com/gpu: 1
            requests:
              memory: 512i
              cpu: 1500m
          volumeMounts:
            - name: mydata
              mountPath: /mnt/myscratch
              command: ["sh", "-c", "sleep infinity"]
      volumes:
        - name: mydata
          emptyDir: {}
          sizeLimit: 10Ti
      nodeSelector:
        nautilus.io/disktype: nvme

```

- Start here:
 - Send Wayne an email - ws@csun.edu

Additional Resources

- Sometimes, researchers just need an unmanaged or managed (by students, supervised by faculty) resource to host public-facing files and applications
 - [Oregon State University Open Source Lab \(OSL\)](#)
- Recently, CSUN was added to the Cloudbank/2i2C JupyterHub resource (this complements CSUN Apporto and SDSU/CSUSB TIDE)
 - [Cal-ICOR JupyterHub Pilot](#)
- Increasingly, Libretexts is moving beyond "texts" and becoming a complete LMS solution, including a JupyterHub resource
 - General System - [LibreTexts](#)
 - Specific Application - [JupyterHub](#)
- Some researchers want to experiment with real Quantum resources
 - [D-Wave LEAP Quantum Launchpad/D-Wave Learn Program \(D-Wave\)](#)
- Many researchers require an AI system that *is* open, transparent, and reproducible (built *top-down*)

- [NSF National Artificial Intelligence Research Resource Pilot \(NAIRR\)](#)
- Some researchers desire an AI system that *is* open, transparent, and reproducible (built *bottom-up*)
 - [Non-Profit Personal AI Lab \(Kwaai\)](#)

Conferences/Fellowships

- There are plenty of zero-cost and low-cost U.S. domestic events for learning about HPC resources at the *Application*-level.

Name	Venue	Cost	Timeframe
Practice & Experience in Advanced Research Computing (PEARC)	varies	mid \$	late July
Science Gateways (SGX3)	varies	\$0 (NSF)	varies
Confab (DOE)	varies	low \$	early April
Institute for Mathematical and Statistical Innovation (IMSI)	varies	\$0 (NSF)	varies
US-RSE Conference (US-RSE)	varies	\$0 (Sloan)	early October
IEEE e-science	varies	mid \$	mid September

- There are plenty of zero-cost and low-cost U.S. domestic events for learning about HPC resources at the *Infrastructure*-level.

Name	Venue	Cost	Timeframe
Research Computing at Smaller Institutions (RCSI)	Swarthmore, PA	\$0 (NSF)	early June
National Research Platform (NRP)	UCSD, CA	\$600	late January
Supercomputing (SC)	St. Louis, Denver, Atlanta	low \$	mid November
Corporation for Networking and Research (CENIC)	varies	low \$	late March
Southern California Linux Expo (SCaLE)	Pasadena, CA	low \$	early March

- And the list of *International* events for learning about HPR resources is growing quickly.

Name	Venue	Cost	Timeframe
CINI HPC Summer School (CINI)	Naples, Italy	N/A	mid June
International Conference on Scalable Scientific Data Management (SSDBM)	varies	N/A	late June

- The following are some of the Fellowships available:

- [ICICLE: Intelligent CI with Computational Learning in the Environment \(ICICLE\)](#)

National Workshops

- There are plenty of *in-person* events for learning about HPR resources.

Name	Venue	Cost	Timeframe
Minority-Serving Cyberinfrastructure Capabilities (MS-CC)	varies	\$0 (NSF)	late May
Open Science Grig (OSG)	U of Wisconsin-Madison, WI	\$0 (NSF)	late June
HPC and Data Science Summer Institute (SDSC)	UCSD, CA	\$350	late July - early August
NERSC International HPC Summer School (NERSC)	varies	\$0 (DOE)	early July
KNIT (FABRIC)	varies	\$0 (NSF)	mid March
INTERSECT Research Software training (INTERSECT)	Princeton/Alabama	\$0 (NSF)	mid July
SHINE Workshop (SHINE)	Iowa City	\$0 fees	mid August

- There are plenty of *virtual* events for learning about HPR resources.

Name	Venue	Cost	Timeframe
OU Supercomputing Center for Education & Research (OU)	virtual, synchronous	\$0 (NSF)	late June
HPC Pathways (NCSA)	virtual, asynchronous	\$0	on-going
Cornell Roadmaps	virtual	\$0	asynchronous, on-going
HPC Carpentry	in-person and virtual, synchronous	\$0	varies
Ecosystem for Research Networking (ERN) Summit	virtual, synchronous only	\$0	late April
Oklahoma University ACI-REF Virtual Residency (Henry Neeman)	virtual	\$0	recurring
(comprehensive, searchable list of resources)	N/A	\$0	varies

Upskilling - Professional Associations/Societies

- Faculty - These HPC resources should be of use to *Faculty* over time.

- [R OpenSci \(ROpenSci\)](#)
- [PyOpenSci \(pyOpenSci\)](#)
- [JuliaCon \(annual Summer conference abstracts, proceedings\)](#)
- [Framework for Open and Reproducible Research Training \(FORRT\)](#)
- [Open Accelerated Computing \(OpenACC\) \(C/C++ optimizations for research, annual Summer conference\)](#)
- [NumFOCUS \(NumFOCUS\) \(open resource software practices\)](#)
- [Consortium for the Advancement of Scientific Software \(CASS\) \(DoE sponsored\)](#)
- [Center for Open-Source Research Software Stewardship and Advancement \(CORSA\)](#)
- [LF AI & Data \(AI Innovation\)](#)
- [US Research Software Sustainability Institute \(URSSI\)](#)
- [Open Molecular Software Foundation \(OMSF\)](#)
- [ZENODO \(open data/scholarship/publication repository - managed by CERN and OpenAIRE\)](#)
- [Software Carpentries \(software engineering\)](#)
- (and check your discipline's pre-conference workshops and related conference themes for HPC events)
- Staff - These HPC resources should be of use to *Staff* over time.
 - [US Research Software Engineering Association \(US-RSE\)](#)
 - [Campus Research Computing Consortium \(CaRCC\)](#)
 - [Campus Champions](#)
 - [OpenOnDemand](#)
 - [Internet2 Research Engagement](#)
 - [Internet2 NET+](#)
 - [EduCause Research Computing and Data Community Group](#)
- Administration - These HPC resources should be of use to *Administration* over time.
 - [Coalition for Academic Scientific Computing \(CASC\)](#)
 - [Research Software Alliance \(ReSA\)](#)
- Sundry - These HPC resources related to *networking* should be of use to various individuals over time.
 - [ES NET \(DOE\)](#)
 - [Globus](#)
 - [The Quilt](#)
 - [Fabric](#)
- Sundry - These miscellaneous open source research-related resources should be of use to various individuals over time.
 - [Research Data Alliance \(RDA\)](#)
 - [Center for Open Science - Open Software Foundation \(OSF\)](#)
 - [Digital Management Plan Tool \(DMPTOOL\)](#)
 - [UC Open Source Program Offices\)](#)
 - [Professional Development for Instructors Interested in Student Participation in Humanitarian Free and Open Source Software \(POSSE\)](#)
 - [HPC Social \(HPC community development\)](#)
- Sundry - These resources related to *domestic* open science/reproducibility should be of use to various individuals over time.

- [Science Philanthropy Alliance](#)
 - [Crossref \(Crossref\)](#)
- Sundry - These resources related to *international* open science/reproducibility should be of use to various individuals over time.
 - [Coalition for Advancing Research Assessment \(CoARA\)](#)
 - [UNESCO Open Science \(UNESCO\)](#)
 - [Global Research Council \(GRC\)](#)
 - [Society of Software Research Engineering \(SocRSE\) \(provides £1,000 for conference travel too\)](#)
 - [Science for Life Laboratory \(SciLifeLab\)](#)
 - [Turing Way \(Turing Way\)](#)
 - [Digital Research Alliance of Canada](#)
 - [International Conference on Scalable Scientific Data Management](#)
 - [Reach the World \(K-12 supercomputing\)](#)
- Sundry - These charitable organizations provide regular funding for HPC-related and scientific software.
 - [Chan/Zuckerberg Initiative \(Essential Open Source Software for Science \(Cycle 6\)\)](#)
 - [Kavli Foundation \(general science\)](#)
 - [Wellcome Trust \(generally, health-related research software\)](#)
 - [Simons Foundation \(generally, math and physical sciences software\)](#)
 - [Sloan Foundation \(generally, emerging technology of any type\)](#)
 - [Code for Science and Society \(generally, open source software\)](#)
- Sundry - Parody.
 - [Journal of Astrological Big Data Ecology \(like the old "Journal of Irreproducible Results"\)](#)